

# Contents

## 1 Some definitions

- overfitting: Given  $H$ ,  $h$  overfits if  $\exists h' \in H$  such that  $h'$  has smaller error over all the instances even though  $h$  has a smaller error over the training examples.

## 2 Lazy vs Eager

- k-NN, locally weighted regression, and case-based reasoning are lazy
- BACKPROP, RBF is eager (why?), ID3 eager
- Lazy algorithms may use query instance  $x_q$  when deciding how to generalize (can represent as a bunch of local functions). Eager methods have already developed what they think is the global function.

## 3 Decision Trees

### 3.1 ID3 Algorithm

- Constructs trees topdown. Greedy algorithm. Hypothesis space of ID3: set of decision trees. Complete space, maintains only a single hypothesis. Uses all training examples at each step (reduced sensitivity to individual error).
  - $A \leftarrow$  best attribute
  - assign  $A$  as decision attribute for Node
  - for each value of  $A$ , create a descendant of node
  - sort training examples to leaves
  - if examples perfectly classified, stop
  - else iterate over leaves
- $Entropy(S) = \sum_{i=1}^c -p_i \lg(p_i)$  ( $p_i$  is proportion of  $S$  belonging to class  $i$ , also base can vary—what would cause us to do that?)
- $Gain(S, A) = Entropy(S) - \sum_{v \in values(A)} \frac{|S_v|}{|S|} Entropy(S_v)$ 
  - $S_v$ : subset of  $S$  for which attribute  $A$  has value  $v$

### 3.2 Inductive Bias of ID3

- prefers shorter trees
- highest info gain attributes

### 3.3 Pruning

- Reduced error (?)
- Rule post-pruning (?)
  - grow the tree
  - convert tree into equivalent set of rules
  - prune (generalize) each rule by removing preconditions that result in improving its estimated accuracy
  - sort pruned rules by estimated accuracy. Consider them in this sequence when classifying subsequent instances.

### 3.4 Adapting Decision Trees to Regression(?)

- splitting criteria: variance
- leaves: average local linear fit

## 4 Regression and Classification

- Least squared error: The objective consists of adjusting the parameters of a model function to best fit a data set. A simple data set consists of  $n$  points (data pairs)  $(x_i, y_i)$ ,  $i = 1, \dots, n$ , where  $x_i$  is an independent variable and  $y_i$  is a dependent variable whose value is found by observation. The model function has the form  $f(x, \beta)$ , where the  $m$  adjustable parameters are held in the vector  $\beta$ . The goal is to find the parameter values for the model which "best" fits the data. The least squares method finds its optimum when the sum,  $S$ , of squared residuals  $S = \sum_{i=1}^n r_i^2$  is a minimum. A residual is defined as the difference between the actual value of the dependent variable and the value predicted by the model.

$r_i = y_i - f(x_i, \beta)$  An example of a model is that of the straight line in two dimensions. Denoting the intercept as  $\beta_0$  and the slope as  $\beta_1$ , the model function is given by  $f(x, \beta) = \beta_0 + \beta_1 x$ .

## 5 Neural Networks

### 5.1 Perceptrons

$$o(x_1 \dots x_n) = \begin{cases} 1, & \text{if } w_0 + w_1 x_1 + \dots + w_n x_n > 0, \\ 0, & \text{otherwise.} \end{cases}$$

where  $w_0, \dots, w_n$  is a real-valued weight. Note that  $w_0$  is a threshold that must be surpassed for the perceptron to output 1. Alternatively:  $o(\vec{x}) = \text{sgn}(\vec{w}\vec{x})$ .  $H = \{\vec{w} | \vec{w} \in \mathbb{R}^{n+1}\}$ .

### 5.2 Perceptron Training Rule vs Delta Rule

- Perceptron training rule: begin with random weights, apply perceptron to each training example, update perceptron weights when it misclassifies. Iterates through training examples repeatedly until it classifies all examples correctly.
  - $w_i \leftarrow w_i + \Delta w_i$

- $\Delta w_i = \eta(t - o)x_i$ ,  $t$ : target output for current training example.  $o$ : output generated for current training example.  $\eta$ : learning rate.
- To converge, Perceptron training rule needs data to be linearly separable( Decision for this hyperplane is  $\vec{w}\vec{x} > 0$ ) and for  $\eta$  to be sufficiently small.
- Delta rule uses *gradient descent*.
  - (?) task of training linear unit (1st stage of a perceptron without the threshold):  $o(\vec{x}) = \vec{w}\vec{x}$
  - training error:  $E(\vec{w}) = \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2$ , where  $D$ : training examples,  $t_d$ : target output for training example  $d$ , and  $o_d$ : output of linear unit for training example  $d$ .
  - Gradient descent finds global minimum of  $E$  by initializing weights, then repeatedly modifying until it hits the global min. Modification: alters in the direction that gives steepest descent.  $\nabla E(\vec{w}) = [\frac{\partial E}{\partial w_0}, \dots, \frac{\partial E}{\partial w_n}]$
  - Training rule for gradient descent:  $w_i \leftarrow w_i + \Delta w_i$   
 $\Delta \vec{w} = -\eta \nabla E(\vec{w})$
  - Training rule can also be written in its component form:  $w_i \leftarrow w_i + \Delta w_i$   
 $\Delta w_i = -\eta \frac{\partial E}{\partial w_i}$
  - Efficient way of finding  $\frac{\partial E}{\partial w_i} = \sum_{d \in D} (t_d - o_d)(-x_{id})$ , where  $x_{id}$  (?) represents single input component  $x_i$  for training example  $d$ .
  - Rewrite:  $\Delta w_i = \eta \sum_{d \in D} (t_d - o_d)(x_{id})$  (true gradient descent)
  - Problems: slow; possibly multiple local minima in error surface (?-I thought error function was smooth, and would always find the global minimum. Example why not?)
  - (?) Stochastic gradient descent:  $\Delta w_i = \eta(t - o)x_i$  (known as delta rule). Error rule:  $E_d(\vec{w}) = \frac{1}{2}(t_d - o_d)^2$  (?-relationship to the other gradient descent? Why don't we need to separate it by  $x_{id}$  anymore? Is this a vector?)
  - Stochastic versus True gradient descent
    - \* true: error summed over all examples before updating weights. stochastic: weights updated upon examining each training example
    - \* summing over multiple examples require more computation per weight update step. But using true gradient, so can use a larger step size
    - \* Stochastic avoids multiple local minima because it uses  $\nabla E_d(\vec{w})$  not  $\nabla E(\vec{w})$
- The cost function for a neural network is non-convex, so it may have multiple minima. Which minimum you find with gradient descent depends on the initialization.

### 5.3 Threshold Unit

Unit for multilayer networks. Want a network that can represent highly nonlinear functions. Need unit whose output is nonlinear, but the output is also differentiable function of its inputs.  $o = \sigma(\vec{w}\vec{x})$  where  $\sigma(y) = \frac{1}{1+e^{-y}}$

### 5.4 BACKPROP

$$E(\vec{w}) = \frac{1}{2} \sum_{d \in D} \sum_{k \in \text{outputs}} (t_{kd} - o_{kd})^2$$

where outputs: set of output units in network,  $t_{kd}$  target,  $o_{kd}$  output associated with  $k^{th}$  output unit and training example  $d$ . (?)

### Algorithm BACKPROP

- until termination condition is met:
- for  $i = 1$  to  $m$  ( $m$  is the number of training examples)
  - set  $a^{(1)} = x^{(i)}$  ( $i^{th}$  training example)
  - Perform forward propagation by computing  $a^{(l)}$  for  $l = 2, \dots, L$  ( $L$  is total number of layers)  $a^{(l)} = \sigma(w^{(l-1)}a^{(l-1)})$  = output of the  $l^{th}$  layer.
  - Using  $y^{(i)}$  compute  $\delta^{(L)} = a^{(L)} - y^{(i)}$  ( $y^{(i)}$  is the target for the  $i^{th}$  training example)
  - Then calculate (??)  $\delta^{(L-1)}$  up until  $\delta^{(2)}$  ( $\delta^{(l)}$  is the "error" of layer  $l$  and
 
$$\delta^{(l)} = w^{(l)}\delta^{(l+1)} \cdot \sigma'(w^{(l)}a^{(l)})$$
  - update  $w^{(l)} = w^{(l)} + \Delta w^{(l)}$  (represents a vector of the weights of layer  $l$ ) where
 
$$\Delta w^{(l)} = \eta \delta^{(l)} \cdot x^{(i)}$$

## 5.5 Momentum

$$\Delta w_n^{(l)} = \eta \delta^{(l)} \cdot x^{(i)} + \alpha w^{(l)}(n-1)$$

where  $n$  is the iteration (adds a momentum  $\alpha$ )

- $E_d(\vec{w}) = \frac{1}{2} \sum_{k \in \text{outputs}} (t_k - o_k)^2$  error on training example  $d$
- How to derive the BACKPROP rule??
- BACKPROP for multi-layer networks may converge only at a local minimum (because error surface for multi-layer networks may contain many different minima).
- Alternative Error Functions?
- Alternative Error Minimization Procedures

**Recurrent Networks** What do I need to know about recurrent networks?

## 5.6 Radial Basis Functions

- $\hat{f}(x) = w_0 + \sum_{u=1}^k w_u \text{Kern}_u(d(x_u, x))$
- Equation can be thought of as training a 2-layer network. First layer computes  $\text{Kern}_u$ , second layer computes a linear combination of these first layer values.
- Kernel is defined such that  $d(x_u, x) \uparrow \implies \text{Kern}_u \downarrow$
- RBF gives global approximation to target function represented by linear combinations of many local kernel functions (smooth linear combination).
- Faster to train than BACKPROP because input and output layer are trained separately.
- RBF is eager: represents global function as a linear combo of multiple local kernel functions. Local approximations RBF creates are not specifically targeted to the query.
- A type of ANN constructed from spatially localized kernel functions. Sort of the 'link' between k-NN and ANN?

## 6 Instance Based Learning

### 6.1 k-NN

- discrete:

$$\hat{f}(x_q) = \operatorname{argmax}_{v \in V} \sum_{i=1}^k \delta(v, f(x_i))$$

where  $\delta(a, b) = 1$  if  $a = b$  and 0 otherwise.

- continuous (for a new value,  $x_q$ ):

$$\hat{f}(x_q) = \frac{\sum_{i=1}^k f(x_i)}{k}$$

- distance-weighted:  $w_i = \frac{1}{d(x_q, x_i)^2}$ . If  $x_q = x_i$  assign  $\hat{f}(x_q) = f(x_i)$  (if more than one, do a majority).

- real valued distance weighted:

$$\hat{f}(x_q) = \frac{\sum_{i=1}^k f(x_i)}{\sum_{i=1}^k w_i}$$

- Inductive Bias of k-NN: assumption that nearest points are most similar
- k-NN is sensitive to having many irrelevant attributes ‘curse of dimensionality’ (can deal with it by ‘stretching the axes’, add a weight to each attribute. Can even get rid of some of the attributes by setting the weight =0)

### 6.2 Locally Weighted Linear Regression

- $f$  approximated near  $x_q$  using  $\hat{f}(\vec{x}) = \vec{w} \cdot \vec{x}$  (is this appropriate notation?)
- Error function using kernel:  $E(x_q) = \frac{1}{2} \sum_{k \in K} (f(x) - \hat{f}(x))^2 \operatorname{Kern}(d(x_q, x))$  where  $K$  is the set of  $k$  closest  $x$  to  $x_q$ .

## 7 Support Vector Machines

Maximal Margin Hyperplanes: if data linearly separable, then  $\exists(\vec{w}, b)$  such that  $\vec{w}^T \vec{x} + b \geq 1 \forall \vec{x}_i \in P$  and  $\vec{w}^T \vec{x} + b \leq -1 \forall \vec{x}_i \in N$  ( $N, P$  are the two classes). Want to minimize  $\vec{w}^T \vec{w}$  subject to constraints of linear separability.

Or, maximize  $\frac{2}{|w|}$  while  $y_i(\vec{w}^T \vec{x}_i + b) \geq 1 \forall i$ . Note  $y_i = \{+1, -1\}$ . Or minimize  $\frac{1}{2}|w|^2$ . This is quadratic programming problem.

$W(\alpha) = \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j x_i^T x_j$ .  $w = \sum_i \alpha_i x_i y_i$ .  $\alpha_i$  mostly 0  $\implies$  only a few of the  $x$ ’s matter.

### 7.1 Kernel Induced Feature Spaces

Map to higher dimensional *feature space*, construct a separating hyperplane.  $X \rightarrow H$  is  $\vec{x} \rightarrow \phi(\vec{x})$ .

Decision function is  $f(\vec{x}) = \operatorname{sgn}(\phi(\vec{x})w^* + b^*)$  (\* means optimal weight and bias)

Kernel function:  $K(\vec{x}\vec{z}) = \phi(\vec{x})^T \phi(\vec{z})$ . If  $K$  exists, we don’t even need to know what  $\phi$  is.

Mercer’s condition:

What if data is not linearly separable? (slack variables?)

Lagrangian?  
Mercer's Theorem?

## 7.2 Relationship between SVMs and Boosting

$H_{trial}(x) = \frac{sgn(\sum_i \alpha_i x_i)}{\sum_i \alpha_i}$ . As we use more and more weak learners, the error stays the same, but the confidence goes up. This equates to having a big margin (big margins tend to avoid overfitting).

## 8 Boosting

Boosting problem: set of weak learners combined to produce a learner with an arbitrary high accuracy.

The original boosting problem asks whether a set of weak learners can be combined to produce a learner with an arbitrary high accuracy. A weak learner is a learner whose performance (at classification or regression) is only slightly better than random guessing. AdaBoost: trains multiple weak classifiers on training data, then combines into single boosted classifier. Weighted sum of weak classifiers with weights dependent on weak classifier accuracy.

$N$  training examples:  $x_i, y_i \in \{-1, +1\}$ . Each example  $i$  has an observation weight  $w_i$  (how important example  $i$  is for our current learning task).

Classifier  $G$ :  $err_G = \sum_{i=1}^N w_i I(y_i \neq G(x_i))$

Using weights:  $err = \frac{\sum_{i=1}^N w_i I(y_i \neq G(x_i))}{\sum_{i=1}^N w_i}$

In this way, our error metric is more sensitive to misclassified examples that have a greater importance weight. Denominator is only for normalization (we want an answer between 0 and  $N$ ).

Boosting: weights are sequentially updated. Algorithm:

- initialize  $w_i = \frac{1}{N}$
- for  $m = 1$  to  $M$ :
  - fit  $G_m(x)$  using  $w_i$ 's
  - compute
 
$$err_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}$$
  - $\alpha_m = \frac{\log(1-err_m)}{err_m}$
  - $w_i \leftarrow w_i \cdot \exp(\alpha_m I(y_i \neq G_m(x_i)))$  for  $i = 1 \dots N$
- $G(x) = \text{sign}[\sum_{m=1}^M \alpha_m G_m(x)]$  In this way, classifiers that have a poor accuracy (high error rate, low  $\alpha_m$ ) are penalized in the final sum.

Question : where are these  $G_m$ 's coming from? Are they pre-set or are they created by the algorithm?

## 9 Computational Learning Theory

### 9.1 Definitions

- $H$ -hypothesis space.  $c \in H$ -true hypothesis.  $h \in H$ -candidate hypothesis.  $S \subseteq H$ -training set.

- Consistent learner: Learner outputs a hypothesis such that  $h(x) = c(x) \forall x \in S$
- Version space:  $VS(S) = \{h \in H : h \text{ consistent wrt to } S\}$  (ie, hypothesis consistent with training examples)
- training error: fraction of training examples misclassified by  $h$ .
- true error: fraction of examples that would be misclassified on sample drawn from  $D$  (distribution over inputs).  $error_D(h) = Pr_{x \sim D}[c(x) \neq h(x)]$
- $C$  is PAC-learnable by learner  $L$  using  $H \iff L$  will output  $h \in H$  (with probability  $1 - \delta$ ) such that  $error_D(h) \leq \varepsilon$  in time and samples polynomial in  $1/\varepsilon, 1/\delta, |H|$ .
- $\varepsilon$ -exhausted version space:  $VS(S)$  exhausted iff  $\forall h \in VS(S) \text{ } error_D(h) \leq \varepsilon$ .

## 9.2 Haussler Theorem

Bounds true error.

Let  $error_D(h_i) > \varepsilon$  for  $i = 1 \dots k$  (some  $h_i$ 's in  $H$ ). How much data do we need to “knock out” all these hypotheses?

$Pr_{x \sim D}[h_i(x) = c(x)] \leq 1 - \varepsilon$  (probability that  $h_i$  matches true concept is low)

$Pr(h_i \text{ consistent with } c \text{ on } m \text{ examples}) \leq (1 - \varepsilon)^m$  (independent).

$Pr(\exists h_i \text{ consistent with } c \text{ on } m \text{ examples}) = k \cdot (1 - \varepsilon)^m \leq |H| \cdot (1 - \varepsilon)^m$

$-\varepsilon \geq \ln(1 - \varepsilon) \implies (1 - \varepsilon)^m \leq \exp(-\varepsilon m)$

Upper bound that VS not  $\varepsilon$ -exhausted after  $m$  samples:  $|H| \cdot \exp(-\varepsilon m)$ .

Want:  $|H| \cdot \exp(-\varepsilon m) \leq \delta$  (solve for  $m$ ).

$m \geq \frac{1}{\varepsilon}(\ln(|H|) + \ln(\frac{1}{\delta}))$

## 9.3 Infinite Hypotheses Spaces

- Examples: linear separators, ANNs, decision trees (continuous inputs)
- $m \geq \frac{1}{\varepsilon}(8VC(H)\lg(\frac{13}{\varepsilon}) + 4\lg(\frac{2}{\delta}))$
- shatter: A set of instances  $S$  is shattered by  $H$  if every possible dichotomy of  $S \exists h \in H$  that is consistent with this dichotomy.
- $VC(H)$  is size of largest finite subset of instance space that can be shattered by  $H$ .
- $C$  PAC-learnable iff VC dimension is finite.

# 10 Bayesian Learning

## 10.1 Equations and Definitions

- $P(h)$  : probability that a hypothesis  $h$  holds
- $P(D)$ : probability that training data  $D$  will be observed
- Bayes' Rule:

$$P(h|D) = \frac{P(D|h)P(h)}{P(D)}$$

- Find most probable  $h \in H$  given  $D$ :

$$h_{map} = \operatorname{argmax}_{h \in H} P(h|D) = \operatorname{argmax}_{h \in H} P(D|h)P(h)$$

- if every  $h \in H$  a priori equally probable:

$$h_{ml} = \operatorname{argmax}_{h \in H} P(D|h)$$

**BRUTE FORCE MAP learning algorithm** Output  $h_{map}$

Let's assume:

- $D$  is noise-free
- Target function  $c \in H$
- all  $h$  (a priori) are equally likely

Then  $P(h) = \frac{1}{|H|}$

$$P(D|h) = \begin{cases} 1, & \text{if } d_i = h(x_i) \forall d_i \in D, \\ 0, & \text{otherwise.} \end{cases}$$

$$P(D) = \frac{|VS_{H,D}|}{|H|}$$

$|VS_{H,D}|$  is the set of hypotheses in  $H$  that are consistent with  $D$ . Consistent learned outputs an  $h$  with zero error over training examples.

Therefore

$$P(h|D) = \begin{cases} \frac{1}{|VS_{H,D}|}, & \text{if } h \text{ consistent with } D \\ 0, & \text{otherwise.} \end{cases}$$

Every consistent hypothesis is a MAP hypothesis (with these assumptions)!

## 10.2 ML and Least-Squared Error

Under certain assumptions any learner that minimizes squared error between the outputs of hypothesis  $h$  and training data will output an ML hypothesis. No idea why. ?? ML hypothesis is the one that minimizes the sum of squared errors over the training data.

## 10.3 Bayes Optimal Classifier

$$P(v_j|D) = \sum_{h_j \in H} P(v_j|h_i)P(h_i|D)$$

(probability that correct classification is  $v_j$ )

$$v_{map} = \operatorname{argmax}_{v_j \in V} P(v_j|D)$$

## 10.4 Bayesian Belief Networks

**Naive Bayes** Classify given attributes:  $v_{map} = \operatorname{argmax}_{v_j \in V} P(v_j|a_1, \dots, a_n)$ . Rewrite using Bayes' rule and use naive assumption that all  $a_i$  are conditionally independent given  $v_j$ .  $v_{NB} = \operatorname{argmax}_{v_j \in V} P(v_j) \prod_i P(a_i|v_j)$ .

Whenever naive assumption is satisfied,  $v_{NB}$  same as MAP classification.



## EM Algorithm

- arbitrary initial hypothesis
- repeatedly calculates expected values of the hidden variables
- recalculates the ML hypothesis

This will converge to local ML hypothesis, along with estimated values for hidden variables (why?)

## 11 Evaluating Hypotheses

## 12 Randomized Optimization

### 12.1 MIMIC

Directly model distribution.

Algorithm:

- generate samples from  $P^{\theta_t}(x)$
- set  $\theta_{t+1}$  to the n'th percentile
- retain only those samples such that  $f(x) \geq \theta_{t+1}$
- estimate  $P^{\theta_{t+1}}(x)$
- repeat!

### 12.2 Simulated Annealing

Algorithm:

- for finite number of iterations:
- sample new point  $x_t$  in  $N(x)$
- Jump to new sample with probability  $P(x, x_t, T)$
- decrease  $T$

$$P(x, x_t, T) = \begin{cases} 1, & \text{if } f(x_t) \geq f(x), \\ \exp(\frac{f(x_t) - f(x)}{T}), & \text{otherwise.} \end{cases}$$

**Genetic Algorithms** WHAT IS??

## 13 Information Theory

**Definitions** We'll use shorthand: Just write  $x$  instead of  $X = x$  for all the possible values that a random event  $X$  could take on. (Am I using the terms correctly?)

- Mutual Information:  $I(X, Y) = H(X) - H(X|Y)$
- Entropy:  $H(A) = -\sum_{s \in A} P(s) \lg(P(s))$
- Joint entropy:  $H(X, Y) = -\sum_{x \in X} \sum_{y \in Y} P(x, y) \lg(P(x, y))$
- Conditional Entropy:  $H(Y|X) = -\sum_{x \in X} \sum_{y \in Y} P(x, y) \lg(P(y|x))$
- If X independent of Y:  $H(Y|X) = H(Y)$  and  $H(Y, X) = H(Y) + H(X)$
- Kullback-Leibler divergence:  $KL(p||q) = -\sum_{x \in X} p(x) \lg(\frac{p(x)}{q(x)})$  for two different distributions  $p, q$ .